



CSE 414 DATABASE
SPRING 2024
HOMEWORK 2

MERT GÜRŞİMŞİR
1901042646

CONTENTS

CONTENTS.....	2
DESIGN	3
SQL FUNCTIONS	5
TABLE SQL FUNCTION	6
FOR LOOP FUNCTION.....	7
FUNCTION WITH VARIABLE.....	8
TRIGGERS	9
TRIGGER WITH “OLD” AND “NEW”	9
TRIGGER WITH “WHEN” AND “IF”	10
TRIGGER WITH “FOR EACH ROW”	12
DROP A TRIGGER & SHOW TRIGGERS	13
TRANSACTIONS	15
FIRST ATOMIC TRANSCATION.....	15
SECOND ATOMIC TRANSACTION.....	16

DESIGN

In this homework, I am going to develop a management system database related to health sector step by step. For the building of database system for health sector operations, I have used PostgreSQL.

There are going to be 7 tables for my management system which are listed below:

- Patients

```
CREATE TABLE Patients (  
    PatientID INT PRIMARY KEY,  
    FirstName VARCHAR(50),  
    LastName VARCHAR(50),  
    DateOfBirth DATE,  
    Gender CHAR(1),  
    Address VARCHAR(100),  
    Phone VARCHAR(15)  
);
```

- Doctors

```
CREATE TABLE Doctors (  
    DoctorID INT PRIMARY KEY,  
    FirstName VARCHAR(50),  
    LastName VARCHAR(50),  
    Specialization VARCHAR(50),  
    Phone VARCHAR(15)  
);
```

- Appointments

```
CREATE TABLE Appointments (  
    AppointmentID INT PRIMARY KEY,  
    PatientID INT,  
    DoctorID INT,  
    AppointmentDate DATE,  
    AppointmentTime TIME,  
    Reason VARCHAR(100),  
    Status VARCHAR(20),  
    FOREIGN KEY (PatientID) REFERENCES Patients(PatientID),  
    FOREIGN KEY (DoctorID) REFERENCES Doctors(DoctorID)  
);
```

- Records

```
CREATE TABLE Records (
    RecordID INT PRIMARY KEY,
    PatientID INT,
    DoctorID INT,
    DateOfRecord DATE,
    Diagnosis VARCHAR(100),
    Treatment VARCHAR(100),
    Prescription VARCHAR(100),
    FOREIGN KEY (PatientID) REFERENCES Patients(PatientID),
    FOREIGN KEY (DoctorID) REFERENCES Doctors(DoctorID)
);
```

- Departments

```
CREATE TABLE Departments (
    DepartmentID INT PRIMARY KEY,
    Name VARCHAR(50),
    Location VARCHAR(100)
);
```

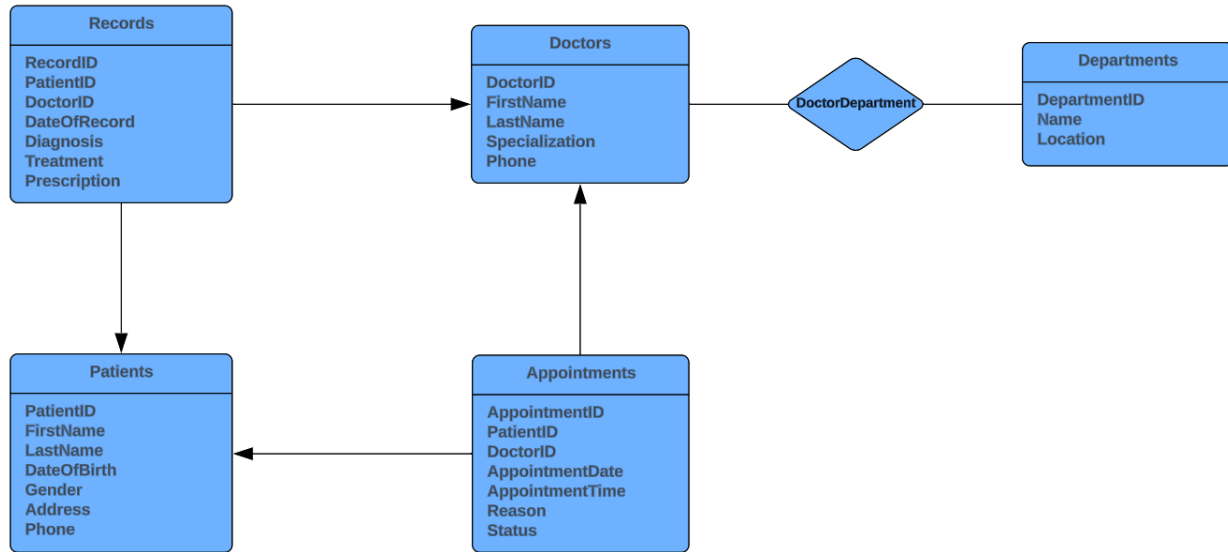
- DoctorDepartment

```
CREATE TABLE DoctorDepartment (
    DoctorID INT,
    DepartmentID INT,
    PRIMARY KEY (DoctorID, DepartmentID),
    FOREIGN KEY (DoctorID) REFERENCES Doctors(DoctorID),
    FOREIGN KEY (DepartmentID) REFERENCES Departments(DepartmentID)
);
```

Below the description of the relationships are listed:

TABLES	CARDINALITY CONSTRAINT	DESCRIPTION
Doctors – Appointments	One-to-Many	A doctor can have many appointments.
Doctors – Records	One-to-Many	A doctor can write many records.
Doctors – Departments	Many-to-Many	A doctor can work in many departments and a department can have many doctors.
Patients – Appointments	One-to-Many	A patient can have many appointments.
Patients – Records	One-to-Many	A patient can have many records.

I am representing my schema diagram below:



SQL FUNCTIONS

First of all, I am adding some values to the tables related to the functions that I am going to write:

```
INSERT INTO Doctors (DoctorID, FirstName, LastName, Specialization, Phone)
VALUES
    (1, 'Sertaç', 'Sever', 'Cardiology', '555-555-5555'),
    (2, 'Agop', 'Kotogyan', 'Dermatology', '555-444-4444'),
    (3, 'Mert', 'Gursimsir', 'Brain Surgeon', '555-333-3333');

INSERT INTO Patients (PatientID, FirstName, LastName, DateOfBirth, Gender, Address, Phone)
VALUES
    (1, 'Stephen', 'King', '1947-09-21', 'M', 'Suite 419 4729 Lemke Plains, East Audria, ME 84366', '532-532-5353'),
    (2, 'Dan', 'Brown', '1964-06-22', 'M', 'Apt. 102 556 Feil Lake, Zoilatown, ID 56072-1825', '532-532-5252');

INSERT INTO Appointments (AppointmentID, PatientID, DoctorID, AppointmentDate, AppointmentTime, Reason, Status)
VALUES
    (1, 1, 3, '2024-06-30', '09:00', 'Check-up', 'Scheduled'),
    (2, 1, 2, '2024-06-30', '10:00', 'Check-up', 'Scheduled'),
    (3, 2, 1, '2024-06-30', '11:00', 'Check-up', 'Scheduled'),
    (4, 2, 3, '2024-06-30', '12:00', 'Check-up', 'Scheduled');
```

These tables look like these now:

Doctors

	doctorid [PK] integer	firstname character varying (50)	lastname character varying (50)	specialization character varying (50)	phone character varying (15)
1	1	Sertaç	Sever	Cardiology	555-555-5555
2	2	Agop	Kotogyan	Dermatology	555-444-4444
3	3	Mert	Gursimsir	Brain Surgeon	555-333-3333

Patients

	patientid [PK] integer	firstname character varying (50)	lastname character varying (50)	dateofbirth date	gender character	address character varying (100)	phone character varying (15)
1	1	Stephen	King	1947-09-21	M	Suite 419 4729 Lemke Plains, East Audria, ME 84366	532-532-5353
2	2	Dan	Brown	1964-06-22	M	Suite 419 4729 Lemke Plains, East Audria, ME 84366	532-532-5353

Appointments

	appointmentid [PK] integer	patientid integer	doctorid integer	appointmentdate date	appointmenttime time without time zone	reason character varying (100)	status character varying (20)
1	1	1	3	2024-06-30	09:00:00	Check-up	Scheduled
2	2	1	2	2024-06-30	10:00:00	Check-up	Scheduled
3	3	2	1	2024-06-30	11:00:00	Check-up	Scheduled
4	4	2	3	2024-06-30	12:00:00	Check-up	Scheduled

Now I am going to write functions needed for the management system one by one.

Table SQL Function

I want to write a table SQL function that gets the ID of a patient and returns the appointments of that patient. I have written a function that is below and called it as you can see:

```
CREATE FUNCTION GetAppointmentsOfAPatient(int)
RETURNS
setof Appointments AS '
SELECT * FROM Appointments WHERE PatientID = $1;
' LANGUAGE SQL;

SELECT * FROM GetAppointmentsOfAPatient(1);
```

The result is this:

	appointmentid integer	patientid integer	doctorid integer	appointmentdate date	appointmenttime time without time zone	reason character varying (100)	status character varying (20)
1	1	1	3	2024-06-30	09:00:00	Check-up	Scheduled
2	2	1	2	2024-06-30	10:00:00	Check-up	Scheduled

For Loop Function

I have written an SQL for loop function that loops through all the appointments and writes down the information about each of the appointments. The function is like this:

```
CREATE OR REPLACE FUNCTION loopAppointments()  
RETURNS VOID AS $$  
    DECLARE  
        appointment RECORD;  
BEGIN  
    FOR appointment IN (SELECT AppointmentDate, DoctorID, PatientID FROM Appointments)  
    LOOP  
        RAISE NOTICE 'There is appointment on % for doctor with ID % by patient %', appointment.AppointmentDate,  
            appointment.DoctorID,  
            appointment.PatientID;  
    END LOOP;  
END;  
$$ LANGUAGE PLPGSQL;
```

I have called it as follows:

```
DO $$  
BEGIN  
    PERFORM loopAppointments();  
END;  
$$;
```

As a result, I got the following:

```
NOTICE: There is appointment on 2024-06-30 for doctor with ID 3 by patient 1  
NOTICE: There is appointment on 2024-06-30 for doctor with ID 2 by patient 1  
NOTICE: There is appointment on 2024-06-30 for doctor with ID 1 by patient 2  
NOTICE: There is appointment on 2024-06-30 for doctor with ID 3 by patient 2
```

Function with Variable

I have written a function that does the following: takes the patient ID as input, use a temporary variable that is total (means the total cost), calculates how many times a patient visits the hospital and multiply by 50 (each appointment costs 50 dollars). Then I have called the function at the bottom.

```
CREATE FUNCTION getTotalCost(patient_ID INT)
RETURNS DECIMAL(10, 2) AS $$
DECLARE
    total DECIMAL(10, 2);
BEGIN
    SELECT COUNT(*) * 50
    INTO total
    FROM Appointments
    WHERE PatientID = patient_ID;

    RETURN total;
END;
$$ LANGUAGE plpgsql;

SELECT getTotalCost(2);
```

The result is this:

	gettotalcost numeric
1	100.00

TRIGGERS

Trigger with “OLD” and “NEW”

In PostgreSQL, there is no explicit definition such as “referencing old row as” and “referencing new row as”. The keywords OLD and NEW can be directly used to access the rows.

Therefore, I have used OLD and NEW in my trigger function.

For this trigger, I have created another table called AppointmentChanges. This table logs the appointments when their status have been changed and when they have been changed.

```
CREATE TABLE IF NOT EXISTS AppointmentChanges (  
    ChangeID SERIAL PRIMARY KEY,  
    AppointmentID INT,  
    OldStatus VARCHAR(20),  
    NewStatus VARCHAR(20),  
    ChangeDate TIMESTAMP DEFAULT CURRENT_TIMESTAMP  
);
```

Trigger goes as follows:

```
CREATE OR REPLACE FUNCTION logAppointmentChange()  
RETURNS TRIGGER AS $$  
BEGIN  
    IF OLD.Status IS DISTINCT FROM NEW.Status THEN  
        INSERT INTO AppointmentChanges(AppointmentID, OldStatus, NewStatus)  
        VALUES (NEW.AppointmentID, OLD.Status, NEW.Status);  
    END IF;  
    RETURN NEW;  
END;  
$$ LANGUAGE plpgsql;  
  
CREATE TRIGGER logChanges  
AFTER UPDATE ON Appointments  
FOR EACH ROW  
EXECUTE FUNCTION logAppointmentChange();
```

I am accessing the old row with OLD and new row with NEW.

Current Appointments table is as follows:

appointmentid [PK] integer	patientid integer	doctorid integer	appointmentdate date	appointmenttime time without time zone	reason character varying (100)	status character varying (20)
2	1	2	2024-06-30	10:00:00	Check-up	Scheduled
3	2	1	2024-06-30	11:00:00	Check-up	Scheduled
4	2	3	2024-06-30	12:00:00	Check-up	Scheduled
1	1	3	2024-06-30	09:00:00	Check-up	Scheduled

Now assume the appointment with ID 1 has been cancelled. In that case this situation is going to be logged into the AppointmentChanges table. The query:

```
UPDATE Appointments SET Status = 'Completed' WHERE AppointmentID = 1;

SELECT * FROM AppointmentChanges;
```

At the end, AppointmentChanges table looks like this:

changeid [PK] integer	appointmentid integer	oldstatus character varying (20)	newstatus character varying (20)	changedate timestamp without time zone
1	1	Scheduled	Completed	2024-06-02 10:29:24.792652

NOTE: In PostgreSQL, triggers are written in the PL/pgSQL language, which requires the use of the 'CREATE FUNCTION statement to define the trigger function. Therefore OLD and NEW keywords are used directly inside this trigger function.

Trigger with “when” and “if”

In this trigger, I am aiming to prevent updating doctors' specializations to a invalid value. I have created a whitelist for the specializations and if a doctor's specialization is going to be updated, this must be one of the specializations at the whitelist. Also the specialization must be different from the old specialization. With this way, trigger is triggered when the update is on specialization.

My trigger is as follows:

```
CREATE OR REPLACE FUNCTION checkSpecializationWhitelist()
RETURNS TRIGGER AS $$
BEGIN
    IF NEW.Specialization NOT IN ('Cardiology', 'Neurology', 'Pediatrics', 'Dermatology', 'Brain Surgeon') THEN
        RAISE EXCEPTION 'Doctor cannot change his/her specialization to given value. % is not in list.', NEW.Specialization;
    END IF;
    RETURN NEW;
END;
$$ LANGUAGE plpgsql;

CREATE TRIGGER doctorUpdate
BEFORE UPDATE ON Doctors
FOR EACH ROW
WHEN (NEW.Specialization IS DISTINCT FROM OLD.Specialization)
EXECUTE FUNCTION checkSpecializationWhitelist();
```

When I try to change the specialization of the doctor to a new value that is not in the whitelist, I expect from program to give an error.

Let's update the doctor's specialization, whose ID is 1, to 'Anesthesiology' which is not in the whitelist.

```
UPDATE Doctors SET Specialization = 'Anesthesiology' WHERE DoctorID = 1;
```

Result is as follows:

```
ERROR: Doctor cannot change his/her specialization to given value. Anesthesiology is not in list.
CONTEXT: PL/pgSQL function checkspecializationwhitelist() line 4 at RAISE
```

But if we try to change the specialization to a value that is in whitelist, update is going to be successful.

```
UPDATE Doctors SET Specialization = 'Neurology' WHERE DoctorID = 1;
SELECT * FROM Doctors;
```

Then we can see specialization is changed:

doctorid [PK] integer	firstname character varying (50)	lastname character varying (50)	specialization character varying (50)	phone character varying (15)
2	Agop	Kotogyan	Dermatology	555-444-4444
3	Mert	Gursimsir	Brain Surgeon	555-333-3333
1	Sertaç	Sever	Neurology	555-555-5555

Trigger with “for each row”

Actually, with the former 2 triggers, we have used the “for each row”. Although in this trigger, the importance of it is going to be mentioned and understood deeply.

In this trigger, I am aiming to do this: Firstly I am going to add a column to the AppointmentCount. Without the trigger, when a new appointment occurs we have to increment the AppointmentCount of the doctor by hand. This is not an efficient approach.

```
ALTER TABLE Doctors ADD COLUMN AppointmentCount INT DEFAULT 0;
```

Now my aim is this: When now appointment is inserted into the Appointments table I am going to increment the related doctor’s AppointmentCount by 1 and in case of deletion of an appointment, I am going to decrease the number of AppointmentCount by 1. Also if the status of the appointment is updated, so it is not “Scheduled”, then we can decrease the AppointmentCount by 1 also.

```
CREATE OR REPLACE FUNCTION updatingAppointmentCount()
RETURNS TRIGGER AS $$
BEGIN
    IF TG_OP = 'INSERT' THEN
        UPDATE Doctors SET AppointmentCount = AppointmentCount + 1 WHERE DoctorID = NEW.DoctorID;
    ELSIF TG_OP = 'DELETE' THEN
        UPDATE Doctors SET AppointmentCount = AppointmentCount - 1 WHERE DoctorID = OLD.DoctorID;
    ELSIF TG_OP = 'UPDATE' THEN
        IF OLD.Status = 'Scheduled' AND NEW.Status <> 'Scheduled' THEN
            UPDATE Doctors SET AppointmentCount = AppointmentCount - 1 WHERE DoctorID = NEW.DoctorID;
        END IF;
    END IF;
    RETURN NEW;
END;
$$ LANGUAGE plpgsql;

CREATE TRIGGER appointmentChange
AFTER INSERT OR DELETE OR UPDATE OF Status ON Appointments
FOR EACH ROW
EXECUTE FUNCTION updatingAppointmentCount();
```

Let’s see how the Doctors and Appointments tables looks for now:

doctorid [PK] integer	firstname character varying (50)	lastname character varying (50)	specialization character varying (50)	phone character varying (15)	appointmentcount integer
1	Sertaç	Sever	Neurology	555-555-5555	1
2	Agop	Kotogyan	Dermatology	555-444-4444	1
3	Mert	Gursimsir	Brain Surgeon	555-333-3333	2

appointmentid [PK] integer	patientid integer	doctorid integer	appointmentdate date	appointmenttime time without time zone	reason character varying (100)	status character varying (20)
2	1	2	2024-06-30	10:00:00	Check-up	Scheduled
3	2	1	2024-06-30	11:00:00	Check-up	Scheduled
4	2	3	2024-06-30	12:00:00	Check-up	Scheduled
1	1	3	2024-06-30	09:00:00	Check-up	Scheduled

Let's cancel the appointments of doctor with ID 3. In this case AppointmentCount of Mert should be 0 because there is no appointment for him:

```
UPDATE Appointments SET Status = 'Canceled' WHERE DoctorID = 3;

SELECT * FROM Doctors;
```

The result is as follows:

doctorid [PK] integer	firstname character varying (50)	lastname character varying (50)	specialization character varying (50)	phone character varying (15)	appointmentcount integer
1	Sertaç	Sever	Neurology	555-555-5555	1
2	Agop	Kotogyan	Dermatology	555-444-4444	1
3	Mert	Gursimsir	Brain Surgeon	555-333-3333	0

Thanks to “for each row” used in the trigger, it is executed for each row. So here 2 lines are updated at appointments and therefore trigger has executed 2 times.

Drop a Trigger & Show Triggers

Firstly, let's see the triggers. I want to see them with name, event object table, event manipulation (update, insert, delete, etc.), and action timing (before or after):

```
SELECT trigger_name, event_manipulation, action_timing, event_object_table FROM information_schema.triggers;
```


Result:

	trigger_name name	event_manipulation character varying	action_timing character varying	event_object_table name
1	logchanges	UPDATE	AFTER	appointments
2	doctorupdate	UPDATE	BEFORE	doctors
3	appointmentchange	INSERT	AFTER	appointments
4	appointmentchange	DELETE	AFTER	appointments
5	appointmentchange	UPDATE	AFTER	appointments

We can also see only the names of the triggers:

```
SELECT DISTINCT(trigger_name) FROM information_schema.triggers;
```


Result:

	trigger_name name	
1	appointmentchange	
2	doctorupdate	
3	logchanges	

Now let's drop some triggers. Firstly, drop the appointmentchange trigger and see the triggers again:

```
DROP TRIGGER IF EXISTS appointmentChange ON Appointments;  
  
SELECT DISTINCT(trigger_name) FROM information_schema.triggers;
```

Result:

	trigger_name name	
1	doctorupdate	
2	logchanges	

TRANSACTIONS

First Atomic Transaction

With this transaction, I am aiming to inserting a new record to Records table and updating the status of an appointment to “Completed”. These both operations should be succeed together or fail together.

```
BEGIN;  
  
INSERT INTO Records (RecordID, PatientID, DoctorID, DateOfRecord, Diagnosis, Treatment, Prescription)  
VALUES (1, 1, 3, '2024-06-30', 'Brain Tumor', 'Surgery', 'Afinitor');  
  
UPDATE Appointments  
SET Status = 'Completed'  
WHERE AppointmentID = 1;  
  
COMMIT;
```

Records table after transaction:

recordid [PK] integer	patientid integer	doctorid integer	dateofrecord date	diagnosis character varying (100)	treatment character varying (100)	prescription character varying (100)
1	1	3	2024-06-30	Brain Tumor	Surgery	Afinitor

Appointments table after transaction:

appointmentid [PK] integer	patientid integer	doctorid integer	appointmentdate date	appointmenttime time without time zone	reason character varying (100)	status character varying (20)
2	1	2	2024-06-30	10:00:00	Check-up	Scheduled
3	2	1	2024-06-30	11:00:00	Check-up	Scheduled
4	2	3	2024-06-30	12:00:00	Check-up	Scheduled
1	1	3	2024-06-30	09:00:00	Check-up	Completed

Second Atomic Transaction

In this transaction, I am aiming to adding a new department and assigning doctors to this department. These two operations should be succeed together or fail together.

```
BEGIN;  
  
INSERT INTO Departments (DepartmentID, Name, Location)  
VALUES (1, 'Neurodermatology', 'Main Building');  
  
INSERT INTO DoctorDepartment (DoctorID, DepartmentID)  
VALUES (1, 1), (2, 1);  
  
COMMIT;
```

Departments table after transaction:

departmentid [PK] integer	name character varying (50)	location character varying (100)
1	Neurodermatology	Main Building

DoctorDepartment table after transaction:

doctorid [PK] integer	departmentid [PK] integer
1	1
2	1